



CIA 1

Part B

Question 1

Installing dependencies and importing libraries

```
In [1]: !pip3 install pandas numpy scikit-learn openpyxl --break-system-packages
```

```
Requirement already satisfied: pandas in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (2.3.1)
Requirement already satisfied: numpy in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (2.3.1)
Requirement already satisfied: scikit-learn in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (1.7.2)
Requirement already satisfied: openpyxl in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (3.1.5)
Requirement already satisfied: python-dateutil>=2.8.2 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from pandas) (2025.2)
Requirement already satisfied: scipy>=1.8.0 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from scikit-learn) (1.16.2)
Requirement already satisfied: joblib>=1.2.0 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: et-xmlfile in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from openpyxl) (2.0.0)
Requirement already satisfied: six>=1.5 in /Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

[notice] A new release of pip is available: 25.1.1 -> 26.1.2

[notice] To update, run: pip3 install --upgrade pip

```
In [2]: import pandas as pd
import numpy as np
```

```
In [3]: data = pd.read_excel("data/QP_Data_road_accident_dataset_with_missing.xlsx")
data.head()
```

Out[3]:

	Country	Year	Month	Day of Week	Time of Day	Urban/Rural	Road Type	Weather Conditions	
0	USA	2002	October	Tuesday	Evening	Rural	Street	Windy	
1	UK	2014	December	Saturday	Evening	Urban	Street	Windy	16
2	USA	2012	July	Sunday	Afternoon	Urban	Highway	Snowy	34
3	UK	2017	May	Saturday	Evening	Urban	Main Road	Clear	48
4	Canada	2002	July	Tuesday	Afternoon	Rural	Highway	Rainy	34

5 rows × 30 columns

```
In [4]: numeric = data.select_dtypes(include=np.number).keys()
print("Numeric data columns: ", numeric)
categorical = data.select_dtypes(include="object").keys()
print("Categorical data columns: ", categorical)
```

```
Numeric data columns: Index(['Year', 'Visibility Level', 'Number of Vehicles Involved',
                             'Speed Limit', 'Driver Alcohol Level', 'Driver Fatigue',
                             'Pedestrians Involved', 'Cyclists Involved', 'Number of Injuries',
                             'Number of Fatalities', 'Emergency Response Time', 'Traffic Volume',
                             'Insurance Claims', 'Medical Cost', 'Economic Loss',
                             'Population Density'],
                             dtype='object')
Categorical data columns: Index(['Country', 'Month', 'Day of Week', 'Time of Day', 'Urban/Rural',
                                 'Road Type', 'Weather Conditions', 'Driver Age Group', 'Driver Gender',
                                 'Vehicle Condition', 'Accident Severity', 'Road Condition',
                                 'Accident Cause', 'Region'],
                                 dtype='object')
```

```
In [5]: print("Data types of all columns: ", data.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 132000 entries, 0 to 131999
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               132000 non-null object
1   Year                                  132000 non-null int64
2   Month                                132000 non-null object
3   Day of Week                           132000 non-null object
4   Time of Day                           132000 non-null object
5   Urban/Rural                           132000 non-null object
6   Road Type                             132000 non-null object
7   Weather Conditions                    132000 non-null object
8   Visibility Level                      125400 non-null float64
9   Number of Vehicles Involved           132000 non-null int64
10  Speed Limit                           132000 non-null int64
11  Driver Age Group                       132000 non-null object
12  Driver Gender                          132000 non-null object
13  Driver Alcohol Level                   104893 non-null float64
14  Driver Fatigue                         132000 non-null int64
15  Vehicle Condition                     132000 non-null object
16  Pedestrians Involved                  132000 non-null int64
17  Cyclists Involved                     132000 non-null int64
18  Accident Severity                     132000 non-null object
19  Number of Injuries                    132000 non-null int64
20  Number of Fatalities                  132000 non-null int64
21  Emergency Response Time                118824 non-null float64
22  Traffic Volume                        125400 non-null float64
23  Road Condition                        132000 non-null object
24  Accident Cause                        132000 non-null object
25  Insurance Claims                      132000 non-null int64
26  Medical Cost                          106979 non-null float64
27  Economic Loss                         118800 non-null float64
28  Region                                132000 non-null object
29  Population Density                    125400 non-null float64
dtypes: float64(7), int64(9), object(14)
memory usage: 30.2+ MB
Data types of all columns:  None

```

All the required information has been printed.

All the columns have appropriate data types.

Question 2a

Descriptive statistics for all the columns

```
In [6]: data.describe()
```

Out[6]:

	Year	Visibility Level	Number of Vehicles Involved	Speed Limit	Driv Alcohol Lev
count	132000.000000	125400.000000	132000.000000	132000.000000	104893.0000
mean	2011.973348	275.104755	2.501227	74.544068	0.1104
std	7.198624	129.946180	1.117272	26.001448	0.0668
min	2000.000000	50.001928	1.000000	30.000000	0.0000
25%	2006.000000	162.422604	2.000000	52.000000	0.0538
50%	2012.000000	274.808977	3.000000	74.000000	0.1076
75%	2018.000000	388.070736	3.000000	97.000000	0.1614
max	2024.000000	499.999646	4.000000	119.000000	0.2499

The 'Economic Loss' column shows the highest variance from the mean with a value of 26642.146034 for its standard deviation.

Question 2b

```
In [7]: print("Medical Cost:")
print("Mean: ", data['Medical Cost'].mean())
print("Median: ", data['Medical Cost'].median())
```

```
Medical Cost:
Mean: 22999.326534801297
Median: 22687.63412
```

The median is lesser than the mean. Hence the distribution of 'Medical Cost' is skewed to the right or positively skewed.

Question 3a

```
In [8]: data.isnull().sum()
```

```
Out[8]: Country          0
        Year             0
        Month            0
        Day of Week      0
        Time of Day      0
        Urban/Rural      0
        Road Type        0
        Weather Conditions 0
        Visibility Level 6600
        Number of Vehicles Involved 0
        Speed Limit      0
        Driver Age Group 0
        Driver Gender    0
        Driver Alcohol Level 27107
        Driver Fatigue   0
        Vehicle Condition 0
        Pedestrians Involved 0
        Cyclists Involved 0
        Accident Severity 0
        Number of Injuries 0
        Number of Fatalities 0
        Emergency Response Time 13176
        Traffic Volume   6600
        Road Condition   0
        Accident Cause   0
        Insurance Claims 0
        Medical Cost     25021
        Economic Loss    13200
        Region           0
        Population Density 6600
        dtype: int64
```

```
In [9]: print("Percentage of missing values: ")
        data.isnull().sum() / data.shape[0] * 100.00
```

Percentage of missing values:

```

Out[9]: Country          0.000000
        Year             0.000000
        Month            0.000000
        Day of Week      0.000000
        Time of Day      0.000000
        Urban/Rural      0.000000
        Road Type        0.000000
        Weather Conditions 0.000000
        Visibility Level  5.000000
        Number of Vehicles Involved 0.000000
        Speed Limit      0.000000
        Driver Age Group  0.000000
        Driver Gender     0.000000
        Driver Alcohol Level 20.535606
        Driver Fatigue    0.000000
        Vehicle Condition 0.000000
        Pedestrians Involved 0.000000
        Cyclists Involved 0.000000
        Accident Severity 0.000000
        Number of Injuries 0.000000
        Number of Fatalities 0.000000
        Emergency Response Time 9.981818
        Traffic Volume    5.000000
        Road Condition    0.000000
        Accident Cause    0.000000
        Insurance Claims  0.000000
        Medical Cost      18.955303
        Economic Loss     10.000000
        Region            0.000000
        Population Density 5.000000
dtype: float64

```

Question 3b

Since 'Visibility Level', 'Population Density' and 'Traffic Volume' have 5 percent missing values, we can use mode imputation.

Question 3c

- Emergency Response Time (9.98%): Falls in the 5-10% range. We will use mean/median/mode substitution.
- Driver Alcohol Level (20.5%) & Medical Cost (18.9%): Have >10% missing values. We will use random sample imputation.

Question 3d

Missingness in 'Driver Alcohol Level' can be:

- MAR (Missing At Random): If the missingness depends on other

observed variables like 'Time of Day' or 'Accident Severity'.


- MNAR (Missing Not At Random): If it depends on the alcohol level itself, like, highly intoxicated drivers refusing the test to avoid legal consequences.

Question 4a

```
In [10]: import matplotlib.pyplot as plt
import seaborn as sns

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.countplot(data=data, x='Weather Conditions', hue='Accident Severity', ax=axes[0])
axes[0].set_title('Accidents by Weather Condition and Severity')
axes[0].tick_params(axis='x', rotation=45)


sns.countplot(data=data, x='Road Type', hue='Accident Severity', ax=axes[1])
axes[1].set_title('Accidents by Road Type and Severity')
plt.tight_layout()
plt.show()
```



Question 4b

```
In [11]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.boxplot(data=data, x='Accident Severity', y='Economic Loss', ax=axes[0])
axes[0].set_title('Economic Loss by Accident Severity')

sns.boxplot(data=data, x='Accident Severity', y='Medical Cost', ax=axes[1])
axes[1].set_title('Medical Cost by Accident Severity')
plt.show()
```



Question 4c

```
In [12]: plt.figure(figsize=(8, 5))
sns.scatterplot(data=data, x='Traffic Volume', y='Economic Loss', hue='Accident Severity')
plt.title('Traffic Volume vs Economic Loss by Accident Severity')
plt.show()
```



Question 4d


Interpretations:

1. Weather and Severity: Severe accidents are more seen in foggy/snowy weather. Minor accidents can be seen as the major contributor.
2. Economic Loss by Severity: The boxplots show that Economic Loss and Medical Cost do not vary much across accident severities.
3. Traffic Volume vs Economic Loss: There is no clear linear relationship between traffic volume and economic loss.

Question 5

```
In [13]: numeric_vars = ['Emergency Response Time', 'Medical Cost', 'Economic Loss', 'Driver Alcohol Level']
corr_matrix = data[numeric_vars].corr()

plt.figure(figsize=(6, 4))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



Interpretation: The correlation matrix shows very weak or near-zero correlations

among these variables, meaning they do not strongly linearly depend on each other.

Why categorical variables shouldn't be directly included: Categorical variables lack an inherent numerical order or magnitude. Mathematical operations used to calculate correlation (like variance and covariance) cannot be meaningfully applied to raw text labels without appropriate encoding.

Question 6a

```
In [14]: cols_to_check = ['Medical Cost', 'Economic Loss', 'Traffic Volume', 'Emergency  
  
for col in cols_to_check:  
    Q1 = data[col].quantile(0.25)  
    Q3 = data[col].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
  
    outliers = data[(data[col] < lower_bound) | (data[col] > upper_bound)]  
    print(f"{col} has {len(outliers)} outliers.")
```

```
Medical Cost has 0 outliers.  
Economic Loss has 0 outliers.  
Traffic Volume has 0 outliers.  
Emergency Response Time has 0 outliers.
```

Question 6b

Examination of Medical Cost Outliers: The maximum value for Medical Cost is approximately 50,000, which can be due to extreme accident cases with fatal injuries. Therefore, these are likely meaningful extreme cases rather than data entry errors.

Question 7

- a. Accident trends over time: Trends typically show fluctuations across years and may peak during certain months (e.g., winter months due to poor weather or holiday seasons with more travel).
- b. Environmental and road conditions: Poor weather (snow/rain) and wet/icy roads increase the likelihood of accidents by reducing visibility and tire traction, acting as significant contributing factors.
- c. Factors influencing severity and impact: Accident severity is often most significantly influenced by 'Speed Limit', 'Driver Alcohol Level', and 'Number of

Vehicles Involved', which directly translate to higher physical impact and subsequent medical/economic costs.

d. Geographic variation: Accident frequencies and impacts vary across regions due to differing population densities, infrastructure quality, and local traffic laws.